

PROJECT: SAMPLE EXAM DATA SCIENTIST ASSOCIATE PRACTICAL



Practical Exam: Supermarket Loyalty

International Essentials is an international supermarket chain.

Shoppers at their supermarkets can sign up for a loyalty program that provides rewards each year to customers based on their spending. The more you spend the bigger the rewards.

The supermarket would like to be able to predict the likely amount customers in the program will spend, so they can estimate the cost of the rewards.

This will help them to predict the likely profit at the end of the year.

Data

The dataset contains records of customers for their last full year of the loyalty program.

Column Name	Criteria
customer_id	Unique identifier for the customer. Missing values are not possible due to the database structure.
spend	Continuous. The total spend of the customer in their last full year. This can be any positive value to two decimal places. Missing values should be replaced with 0.
first_month	Continuous. The amount spent by the customer in their first month of the year. This can be any positive value, rounded to two decimal places. Missing values should be replaced with 0.
items_in_first_month	Discrete. The number of items purchased in the first month. Any integer value greater than or equal to zero. Missing values should be replaced by 0.
region	Nominal. The geographic region that the customer is based in. One of four values Americas, Asia/Pacific, Europe, Middle East/Africa. Missing values should be replaced with "Unknown".
loyalty_years	Ordinal. The number of years the customer has been a part of the loyalty program. One of five ordered categories, '0-1', '1-3', '3-5', '5-10', '10+'. Missing values should be replaced with '0-1'.
joining_month	Nominal. The month the customer joined the loyalty program. One of 12 values "Jan",

Column Name	Criteria
	"Feb", "Mar", "Apr", etc. Missing values should be replaced with "Unknown".
promotion	Nominal. Did the customer join the loyalty program as part of a promotion? Either 'Yes' or 'No'. Missing values should be replaced with 'No'.

Task 1

Before you fit any models, you will need to make sure the data is clean.

The table below shows what the data should look like.

Create a cleaned version of the dataframe.

- You should start with the data in the file "loyalty.csv".
- Your output should be a dataframe named `clean_data`.
- All column names and values should match the table below.

Column Name	Criteria
customer_id	Unique identifier for the customer. Missing values are not possible due to the database structure.
spend	Continuous. The total spend of the customer in their last full year. This can be any positive value to two decimal places. Missing values should be replaced with 0.
first_month	Continuous. The amount spent by the customer in their first month of the year. This can be any positive value, rounded to two decimal places. Missing values should be replaced with 0.
items_in_first_month	Discrete. The number of items purchased in the first month. Any integer value greater than or equal to zero. Missing values should be replaced by 0.
region	Nominal. The geographic region that the customer is based in. One of four values Americas, Asia/Pacific, Europe, Middle East/Africa. Missing values should be replaced with "Unknown".
loyalty_years	Ordinal. The number of years the customer has been a part of the loyalty program. One of five ordered categories, '0-1', '1-3', '3-5', '5-10', '10+'. Missing values should be replaced with '0-1'.
joining_month	Nominal. The month the customer joined the loyalty program. One of 12 values "Jan", "Feb", "Mar", "Apr", etc. Missing values should be replaced with "Unknown".
promotion	Nominal. Did the customer join the loyalty program as part of a promotion? Either

Column Name**Criteria**

'Yes' or 'No'.

Missing values should be replaced with 'No'.

```
import pandas as pd

# Load the data
df = pd.read_csv("loyalty.csv")

# Cleaning the DataFrame
# 1. Ensure 'customer_id' has no missing values
df['customer_id'] = df['customer_id'].dropna()

# 2. Clean 'spend'
df['spend'] = df['spend'].fillna(0).round(2)
df['spend'] = df['spend'] # Ensure all values are non-negative

# 3. Clean 'first_month'
df['first_month'] = df['first_month'].fillna(0)
df['first_month'] = df['first_month']

# 4. Clean 'items_in_first_month'
df['items_in_first_month'] = df['items_in_first_month'].fillna(0).astype(int)
df['items_in_first_month'] = df['items_in_first_month']

# 5. Clean 'region'
df['region'] = df['region'].fillna("Unknown")
valid_regions = ["Americas", "Asia/Pacific", "Europe", "Middle East/Africa"]
df['region'] = df['region'].where(df['region'].isin(valid_regions), "Unknown")

# 6. Clean 'loyalty_years'
df['loyalty_years'] = df['loyalty_years'].fillna("0-1")
valid_loyalty_years = ['0-1', '1-3', '3-5', '5-10', '10+']
df['loyalty_years'] =
df['loyalty_years'].where(df['loyalty_years'].isin(valid_loyalty_years), "0-1")

# 7. Clean 'joining_month'
df['joining_month'] = df['joining_month'].fillna("Unknown")
valid_months = ["Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep",
"Oct", "Nov", "Dec"]
df['joining_month'] =
df['joining_month'].where(df['joining_month'].isin(valid_months), "Unknown")

# 8. Clean 'promotion'
df['promotion'] = df['promotion'].fillna("No")
df['promotion'] = df['promotion'].where(df['promotion'].isin(['Yes', 'No']), "No")

# Create the cleaned DataFrame
clean_data = df.copy()

# Convert data types safely
```

```

# 1. customer_id: Ensure it's a string
clean_data['customer_id'] = clean_data['customer_id'].astype(str)

# 2. spend: Convert to float, replacing non-numeric with NaN, then fill NaN with 0
clean_data['spend'] = pd.to_numeric(clean_data['spend'],
errors='coerce').fillna(0).round(2)

# 3. first_month: Convert to float, replacing non-numeric with NaN, then fill NaN
with 0
clean_data['first_month'] = pd.to_numeric(clean_data['first_month'],
errors='coerce').fillna(0).round(2)

# 4. items_in_first_month: Convert to integer, replacing non-numeric with NaN, then
fill NaN with 0
clean_data['items_in_first_month'] =
pd.to_numeric(clean_data['items_in_first_month'],
errors='coerce').fillna(0).astype(int)

# 5. region: Ensure it's a category
clean_data['region'] = clean_data['region'].astype('category')

# 6. loyalty_years: Ensure it's an ordered categorical type
loyalty_years_categories = ['0-1', '1-3', '3-5', '5-10', '10+']
clean_data['loyalty_years'] = pd.Categorical(clean_data['loyalty_years'],
categories=loyalty_years_categories, ordered=True)

# 7. joining_month: Ensure it's a category
clean_data['joining_month'] = clean_data['joining_month'].astype('category')

# 8. promotion: Ensure it's a category
clean_data['promotion'] = clean_data['promotion'].astype('category')

# Display the cleaned DataFrame
print(clean_data)

```

	customer_id	spend	first_month	...	loyalty_years	joining_month	promotion
0	1	132.68	15.30	...	5-10	Nov	No
1	2	106.45	16.20	...	0-1	Feb	Yes
2	3	123.16	25.26	...	10+	Dec	Yes
3	4	130.60	24.74	...	3-5	Apr	No
4	5	130.41	25.59	...	3-5	Apr	Yes
...
1241	1242	126.01	15.97	...	1-3	Unknown	Yes
1242	1243	111.85	21.22	...	0-1	Jan	No
1243	1244	132.51	16.23	...	5-10	Aug	No
1244	1245	131.57	21.21	...	1-3	Unknown	No
1245	1246	135.85	18.99	...	5-10	Unknown	No

[1246 rows x 8 columns]

Task 2

The team at International Essentials have told you that they have always believed that the number of years in the loyalty scheme is the biggest driver of spend.

Producing a table showing the difference in the average spend by number of years in the loyalty programme along with the variance to investigate this question for the team.

- You should start with the data in the file 'loyalty.csv'.
- Your output should be a data frame named `spend_by_years`.
- It should include the three columns `loyalty_years`, `avg_spend`, `var_spend`.
- Your answers should be rounded to 2 decimal places.

Use this cell to write your code for Task 2

```
spend_by_years = df.groupby('loyalty_years')['spend'].agg(  
    avg_spend='mean',  
    var_spend='var'  
) .reset_index()
```

Round the results to 2 decimal places

```
spend_by_years['avg_spend'] = spend_by_years['avg_spend'].round(2)  
spend_by_years['var_spend'] = spend_by_years['var_spend'].round(2)
```

Display the resulting DataFrame

```
print(spend_by_years)
```

	loyalty_years	avg_spend	var_spend
0	0-1	110.56	9.30
1	1-3	129.31	9.65
2	10+	117.41	16.72
3	3-5	124.55	11.09
4	5-10	135.15	14.10

Task 3

Fit a baseline model to predict the spend over the year for each customer.

1. Fit your model using the data contained in “train.csv”
2. Use “test.csv” to predict new values based on your model. You must return a dataframe named `base_result`, that includes `customer_id` and `spend`. The `spend` column must be your predicted values.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

def clean_data(df):
    # Handle missing values
    df.fillna(method='ffill', inplace=True) # Example: forward fill for missing
    values

    # Encode categorical variables (example for one-hot encoding)
    df = pd.get_dummies(df, drop_first=True)

    # Remove duplicates
    df.drop_duplicates(inplace=True)

    return df

# Load the training data
train_data = pd.read_csv("train.csv")
train_data = clean_data(train_data) # Clean the training data

# Prepare the features and target variable
X_train = train_data.drop(columns=['customer_id', 'spend']) # Drop customer_id and
target variable
y_train = train_data['spend']

# Fit the baseline model
baseline_model = LinearRegression()
baseline_model.fit(X_train, y_train)

# Load the test data
test_data = pd.read_csv("test.csv")
test_data = clean_data(test_data) # Clean the test data

# Prepare the test features
X_test = test_data.drop(columns=['customer_id'])

# Predict spend using the baseline model
predicted_spend = baseline_model.predict(X_test)

# Create the result DataFrame
base_result = pd.DataFrame({
    'customer_id': test_data['customer_id'],
    'spend': predicted_spend
})

# Display the result
print(base_result)
```

	customer_id	spend
0	5	140.699644
1	7	148.730919
2	16	140.810384
3	17	150.649670
4	19	153.628831
..
245	1216	134.869770
246	1225	148.542583
247	1231	136.969013
248	1242	129.885730
249	1243	148.518401

[250 rows x 2 columns]

Task 4

Fit a comparison model to predict the spend over the year for each customer.

1. Fit your model using the data contained in “train.csv”
2. Use “test.csv” to predict new values based on your model. You must return a dataframe named `compare_result`, that includes `customer_id` and `spend`. The `spend` column must be your predicted values.

```
from sklearn.ensemble import RandomForestRegressor

# Fit the comparison model
comparison_model = RandomForestRegressor(n_estimators=100, random_state=42)
comparison_model.fit(X_train, y_train)

# Predict spend using the comparison model
predicted_spend_compare = comparison_model.predict(X_test)

# Create the result DataFrame
compare_result = pd.DataFrame({
    'customer_id': test_data['customer_id'],
    'spend': predicted_spend_compare
})

# Display the result
print(compare_result)
```

	customer_id	spend
0	5	140.5627
1	7	148.7941
2	16	141.0780
3	17	150.8226
4	19	153.6214
..
245	1216	134.8089
246	1225	148.5374
247	1231	136.9832
248	1242	129.7062
249	1243	148.5606

[250 rows x 2 columns]

 No source selected